

Grounded, Not Automated: A Verification-First LLM Career System

Kayvan Nejabati Zenouz
ZenGrowth (ZENOUZ.ai)
zenouz.ai@gmail.com

August 2026

Working paper v0.4 — working system and prospective evaluation; comparative measurements pending

Abstract

AI career tools promise more interviews through resume optimization or unattended application volume, usually without audited outcome denominators. We present ZenGrowth, a working, audience-of-one, local-first system used in a live senior AI/Data-Science search. One repository connects compliant role discovery, inspectable scoring, evidence-grounded application materials, interview preparation, offer evaluation, and lifecycle records. Its design constrains generation to a reviewed evidence bank, rejects selected unallowlisted numbers and named entities, records instrumented model calls, and requires operator approval for every outbound artifact. The implementation and regression tests establish capability-level evidence: these workflow and control mechanisms exist and execute. They do not establish comparative speed, sentence-level factual entailment, or improved callback rates. We therefore separate demonstrated functionality from process and outcome claims, and prospectively specify audits of grounding, score stability, review and elapsed time, cost, and downstream outcomes. The current contribution is practical as well as methodological: a coherent, operator-controlled application workflow, a falsifiable evaluation protocol, and a claim-audit checklist. Our working hypothesis is deliberately modest: supervised drafting can be useful when errors remain reviewable, while autonomous submission moves model errors into an irreversible, reputation-bearing action.

1 Introduction

Between 2023 and 2026 the AI job-search category split into two promises. The first is *optimization*: score your resume against a job description and close the keyword gap (Jobscan’s Match Rate, Teal’s Match Score) [17, 23]. The second is *automation*: agents that discover postings and submit applications with minimal human involvement, at volumes of hundreds per week [2, 16]. Both promises are typically supported by testimonials and vendor-reported statistics rather than denominator-complete outcome audits.

ZenGrowth is a system built inside this category but against its grain. It automates everything *except* the send: discovery, explainable scoring, evidence-grounded document generation, interview preparation, and offer evaluation — with a verification path that makes several high-risk forms of fabrication mechanically detectable, and a human-approval boundary that makes every outbound artifact the operator’s deliberate act. It is also an audience-of-one deployment:

built by its operator for a live senior AI/Data-Science search, with workflow events and instrumented model costs logged.

That combination — category-typical capability, category-atypical verification, and an inspectable audit trail — is what makes an honest field note possible. This paper asks three questions. (1) What useful functionality does the working repository already establish (§3)? (2) Which claims about smoothness, speed, correctness, and outcomes still need measurement (§3)? (3) Where do LLMs genuinely add value in the job search, given the evidence behind category claims (§4–§7)?

2 Relation to prior work

Candidate-side LLM systems are no longer absent from the literature. ResumeFlow generates job-specific resumes from a source resume and proposes task-specific alignment and hallucination metrics [29]. Resume Tailor is the closest *longitudinal case study*: a career vault, multi-source retrieval, provenance tracking, guardrails, and a conditional review loop across nine job descriptions, with ATS-style fit-score deltas [1]. Grounded Optimization is the closest *method* paper on rewrite defenses: a five-layer stack (temporal validation, deterministic contamination detection, structural invariants, prompt-level grounding, and an evaluator agent) with ablations on 25 synthetic resumes, reducing detected hallucinations from 2.48–5.36 per resume in undefended baselines to 0.04–0.24 [15]. These results already occupy the “add grounding gates” slot. They do not audit a live search, measure operator correction burden, or report denominator-complete downstream outcomes.

Synapse takes the opposite design bet: evolutionary, LLM-guided resume mutation to raise recommender and screening scores [12]. That is a useful foil. ZenGrowth treats score-chasing as a failure mode (§6), not as an objective.

ConFit v3 studies the adjacent matching problem with multi-pass LLM re-ranking, listwise training, data cleaning, and distillation on real-world person–job fit datasets [26]. It strengthens benchmark evidence for matching and ranking; it does not establish that an integrated candidate workflow is faster, more accurate, or more effective.

Attune shows a newer human–AI direction for text-record scoring: users can inspect and deterministically edit criteria and rules derived from pairwise comparisons [10]. Its formative study ($n = 12$), technical evaluation, and domain-expert study ($n = 8$) make rationale-bearing scoring alone a weak novelty claim. ZenGrowth does not implement Attune’s interactive criteria steering; its narrower opportunity is to measure score stability, operator disagreement, and correction burden in a live candidate workflow.

ZenGrowth’s remaining proposed delta is therefore *operational*: a live, local-first deployment that connects reviewed evidence, deterministic write-time risk gates, cost and latency telemetry, an application-level action log, and a hard approval boundary between reversible assistance and reputation-bearing send. Unlike an ATS-style fit score, its planned evaluation treats unsupported or embellished claims after review as the primary failure outcome, and reports operator time as part of the return rather than a free safety layer. That delta is a hypothesis until the audits in §3 are complete.

In the broader agent literature, ZenGrowth is not a new general-purpose agent architecture or benchmark. AgencyBench evaluates long-horizon capability over 138 tasks and 32 real-world scenarios [18]. Auditable Agents already formalizes action recoverability, lifecycle coverage, policy checkability, responsibility attribution, and evidence integrity [20]. ZenGrowth should

be read as a domain-specific implementation and evaluation opportunity, not as the origin of agent auditability. Interview research with 17 experienced developers identifies a priori control, co-planning, real-time monitoring, and post hoc review as distinct forms of oversight [11]. ZenGrowth currently emphasizes a priori gates and post hoc operator review; it does not yet evaluate the other forms or show that its review is sufficient. Recent work on audit-budget allocation further shows why self-reported model confidence is an unsafe basis for selective review when confidence is miscalibrated or errors are correlated [27].

Employer-side work provides important context but not evidence that candidate-side tailoring improves outcomes. In two recruitment-platform field experiments, shortlisted candidates assessed with AI interview information passed a blind final human interview 17.5–20 percentage points more often than candidates shortlisted from resumes alone; however, 75% of invitees did not complete the AI interview, shifting screening cost to applicants [3]. This is evidence for a different intervention, not for resume rewriting or application agents. Classical correspondence experiments manipulate perceived identity in otherwise similar applications [8]. More recent audits find demographic sensitivity in embedding retrieval and GPT-based resume assessment [5, 24]; a 2026 audit across newer models reports null or reversed gaps in several settings, showing that direction and magnitude depend on model vintage and protocol [13]. Other 2026 counterfactual audits find disparities from subtle sociocultural markers and gender effects in Japanese-format resumes, reinforcing that anonymization, prompting, language, and task design materially affect an audit result [14, 22]. Nor does a human-in-the-loop label guarantee correction: in a controlled resume-screening experiment ($N = 528$), participants followed simulated biased AI recommendations up to 90% of the time, including in some conditions where they judged the recommendations low quality or unimportant [25]. The relevant conclusion is not that all screeners share one stable bias, but that candidate materials enter a model-mediated environment whose behavior requires versioned, task-specific auditing, and that human approval must itself be evaluated rather than assumed sufficient.

3 Capability evidence: the working ZenGrowth deployment

3.1 The system under test

ZenGrowth is a local-first FastAPI/React application with a SQLite store, operated since May 2026 for a live senior AI/DS search. Its pipeline: nightly compliant discovery over public ATS feeds (Greenhouse/Lever JSON; no scraping of job boards) with a dead-man’s-switch heartbeat; an LLM-free relevance precheck; single-call strict-JSON scoring that records a per-dimension rationale and a *priority score* — a weighted sum of observable fit dimensions in which the model’s uncalibrated success-probability guess is demoted to a tie-breaking band and application effort is a separate cost axis; and evidence-grounded generation of CVs, cover letters, and application answers. Interview packs, transcript debriefs, offer evaluation, and onboarding/departure planning use the same pattern: research and drafting with provenance, never dispatched by the system.

The verification path is the load-bearing component (Figure 1). Source documents are parsed into claims with provenance; claims auto-verify only when the extractor’s confidence clears a configured bar, the cited span exists in the source, and the claim’s high-risk tokens match that span, while the rest queue for operator review. Generation receives only verified claims, or, when the bank holds none, a hand-authored evidence file the operator maintains. Write-time gates reject numbers and a conservative subset of named entities found in neither the evidence bank




Figure 1: Verification-first pipeline. Evidence is verified before generation. Deterministic gates catch selected token classes. Only the operator may cross into an outbound action. The control plane records spend and workflow transitions so later aggregates are diagnosable.

nor the job context. These allowlists catch important fabrication classes — invented figures, and uncited names in CamelCase or acronym form — but they are token allowlists, not entailment checks. Single-capitalised names (most employer names) are not detected, an allowed token may be reused in an unsupported context, and auto-verification itself checks hard facts against the cited span rather than meaning. Whether each sentence is entailed by its evidence, and how often the auto-verified pool is a faithful paraphrase of its source, are both subjects of the planned claim-level audit. Alias resolution merges entity surface variants. A deterministic, zero-LLM-cost JD match report (term coverage, quantified-line count, cliché detection) accompanies every material as a signal, not a target. Drafts carry a never-sent banner until the operator acts.

Instrumented model paths record cost and latency; a configured daily spend ceiling fails closed; and meaningful workflow transitions land in an audit log streamed to the dashboard. The continuously integrated test suite includes deterministic eval gates: hard-fact faithfulness, retrieval recall, entity-resolution pairwise F1, claim-versus-span distortion checks, and a scoring regression suite (rewording stability bounds, hand-ranked monotonicity, anchor drift alarms) — all CI-enforced with zero LLM spend. These run the gate and aggregation logic over small hand-authored fixture sets, each with a negative case proving the check is not vacuous; they are regression protection for the mechanisms, not measurements of generated output or of live scorer stability, which remain the prospective quantities in Table 2.

Scoring is *model-aware*. Older Anthropic models may receive an explicit `temperature=0` as a reproducibility request. Current production defaults to Claude Sonnet 5, which rejects non-default sampling parameters; the client omits temperature rather than send an invalid request. Neither choice guarantees identical provider outputs. Repeated-score stability is therefore a measured quantity, not an assumed property of the API. Sonnet 5’s new tokenizer also prevents direct reuse of earlier token and cost baselines; model, tokenizer, pricing date, and token counts must be frozen together [4].

3.2 Capability evidence and the utility boundary

The repository is more than an architectural proposal or a collection of prompts: it runs an integrated workflow whose outputs remain inspectable by its operator. Table 1 distinguishes that demonstrated functionality from the comparative claims that begin with measurement.

Capability	Repository evidence	Defensible current claim
Find and prioritize	Feed adapters, deterministic precheck, rationale-bearing score	Working intake and prioritization; not proven relevance or ranking benefit
Prepare materials	Reviewed claim bank, provenance, numeric/entity gates	Reviewable, evidence-bounded drafting; not complete semantic correctness
Continue the process	Interview, offer, onboarding, and departure modules	One evidence trail across stages; not measured effort reduction
Control and diagnose	Never-send boundary, telemetry, spend ceiling, audit events	Operator-controlled and inspectable operation; not proof every error is caught

Table 1: Capability evidence ladder. The working repository supports the claims in the final column now; speed, smoothness, correctness rates, and downstream outcomes require the prospective evaluation below.

This is the paper’s present utility claim: ZenGrowth supplies useful functionality for organizing and executing a job-application workflow while keeping external action under human control. Whether it makes that workflow smoother or faster, how often its drafts remain correct after review, and whether it improves application outcomes are separate empirical questions.

3.3 Prospective evaluation and current status

Table 2 is a claims scorecard, not a results table. The repository establishes that the listed mechanisms exist and are tested; it does not establish their field rates or downstream value. Measurements will be frozen as dated aggregate artifacts using the public protocol [28] before numerical claims enter the abstract.

Claim	Current verdict	Required measurement
C3/C7	Partially verified: gates exist; semantic fidelity unmeasured	Precision-targeted, stratified pilot of ≥ 75 claims; pre/post-review grounded, unsupported, and embellished rates with exact intervals
C4	Unmeasured: “about two minutes” withdrawn pending timing	All eligible runs in a frozen window (target ≥ 20); wall-clock split into wait, model, review, and correction
C5	Unmeasured: per-call telemetry exists	Per-application cost distribution with operation inclusion rules and currency conversion date
C2/C8	Partially verified: rationale exists; sampling is model-aware	≥ 10 frozen postings with five repeats per model/version; score and rank variance
C1	Unmeasured: discovery pipeline exists	Archive audit with relevance rubric, sample size, and relevant-role loss rate
C6	Partially verified: local store and documented model egress	Enumerated egress audit: payload class, recipient, trigger, retention, and disclosure
C2b	Unmeasured: usefulness of rank unknown	Rank association with pre-specified pursue decisions; report n and dependence limitations
—	Unmeasured: callback/interview effect	Tracked campaign with complete application denominator and outcome follow-up

Table 2: Claim status and prospective measurement plan for v1.

3.4 What the working repository already demonstrates

Some findings do not wait for the quantitative pass because they are architectural facts with test witnesses. The grounding gates fail generation when a detectable, unallowlisted figure or entity appears, and the distortion check demotes extracted claims whose numbers do not

match their cited span. CV page-fit rewrites face the same discipline plus structure-preserving checks (no new ungrounded figures/entities/words, no change to bullet or section counts, and no added content-hiding constructs) before the fitted `.tex` is kept. Scoring records a rationale and applies sampling parameters only when the named model accepts them. A configured daily ceiling blocks further model calls once recorded spend reaches the cap. These are properties of the implementation, not estimates of end-to-end factual accuracy, scoring stability, or user value. The system is also *diagnosable*: logged workflow transitions and instrumented model costs make the planned aggregate analyses possible. That auditability supported the diagnosis of ranking-formula and lifecycle defects during development, while the completeness of the log remains something to test rather than assume. The evidence ladder is therefore explicit: working capability is established by implementation and tests; process benefit is currently qualitative; effect size, error rate, time saving, and downstream value remain unmeasured.

3.5 Limitations of our evidence

A single operator, a single live search, one configuration lineage, no control arm (no paired untailored applications), and small n throughout; the operator is also the system’s author, so pursue-decisions are not independent of the scoring being evaluated. Downstream outcomes confound with market conditions, seniority, and field. The grounding audit samples materials generated for real applications and therefore cannot be published raw — only aggregate, de-identified rates. Grounded Optimization already reports rewrite-defense rates on synthetic resumes [15]; this paper cannot import those rates as field performance. These limitations bound the claims: our data will support “measured properties of one audited deployment, with diagnosed mechanisms,” not effect sizes for the category.

4 The AI-career-tool claims landscape

Qualitative synthesis from primary product pages accessed in 2026-07. Product-defined match scores and vendor outcome claims are treated as marketing evidence, not causal estimates. Academic neighbors were rechecked through 2026-08-26.

4.1 Optimization tools: scores without outcome audits

Jobscan recommends a product-defined Match Rate of at least 75%; Teal exposes a 0–100% resume-to-job score and emphasizes keyword and skill alignment [17, 23]. Both publish outcome-oriented claims, but the reviewed pages do not disclose assignment mechanisms, denominator-complete cohorts, attrition, or independent audits. Their scores may be useful editing signals, but they are not calibrated callback probabilities. Classical correspondence studies supply controlled callback evidence for employer discrimination [8]; they do not identify the causal effect of commercial match-score optimization for the same candidate. Systems that mutate a resume to raise a screener’s score [12] inherit this conflation unless they also report unsupported-claim rates and downstream outcomes.

4.2 Automation tools: volume as value

JobCopilot advertises up to 50 applications per day and “10X more” interviews; its public page supports the product capability and the existence of the claim, not its causal validity [16]. AIHawk documents an open-source automated application agent; the original repository was

archived in May 2026 [2]. Volume is an activity metric. A useful evaluation must also report complete outcome denominators, review time, false submissions, and corrections. Claims about employer detection or platform-wide countermeasures are excluded here because this review did not identify a sufficiently direct primary evidence base.

4.3 The screening environment

Generated applications can enter algorithmic or LLM-mediated screening. Published audits demonstrate demographic sensitivity, but disagree on its direction across models and protocols [5, 13, 24]. For a candidate-side tool this cuts two ways: keyword optimization may exploit model-specific ranking behavior rather than communicate fit, and a fabricated or embellished claim may survive early screening only to create greater risk at interview or verification stages.

5 Cross-domain evidence

No reviewed experiment establishes that tailoring applications with an LLM improves callback rates. Employer-side AI interviewing now has field evidence [3], but it is a different intervention with a different cost allocation. The following cross-domain studies motivate evaluation choices and design hypotheses only; their effect sizes do not transfer to job search.

Bounded assistance. A staggered rollout across 5,179 support agents increased resolutions per hour by 14% on average and 34% for novice and lower-skilled workers, without reducing average customer satisfaction [9]. A separate Copilot experiment found 55.8% faster completion on one specified JavaScript task [21]. These results show that bounded assistance can help in some settings; they do not identify the value of resume tailoring. ZenGrowth’s operator-review boundary is an evidence-informed design choice to be tested, not a result imported from those studies.

Perception is not elapsed time. An early-2025 METR randomized trial of 16 experienced open-source developers found a 19% slowdown while participants estimated after the study that AI had made them 20% faster [6]. METR’s later experiment encountered substantial task and participant selection effects as agentic tools became harder to withhold [7]. The appropriate lesson is methodological: self-reported productivity and task-level timing answer different questions, and both require explicit sampling assumptions.

Supervision has a cost. Agent-usability work frames human supervision time as part of the return-on-investment calculation rather than a free safety layer [19]. This matters in both directions: review can prevent an irreversible submission error, but excessive review can erase any time saving. ZenGrowth therefore needs to measure operator review time and correction burden alongside model latency and spend.

6 A taxonomy of claim failure modes for AI career tools

Grounded Optimization defines *document-rewrite* failures: temporal injection, cross-domain contamination, structural mutation, and content fabrication [15]. Table 3 defines *category-claim* failures: how tools describe their value. Both are needed. A system can pass a token-level rewrite check and still sell volume as if it were interviews.

Failure mode	Mechanism	Category example
Survivorship testimonial	Winners self-report; denominators absent	“I got 5 interviews in a week”
ATS folklore	Unverifiable claims about parser behavior drive product design	“Robots reject resumes without keyword X”
Score–outcome conflation	Keyword coverage quoted as if it were callbacks	Match-score marketing
Volume-as-value	Applications sent quoted; interviews per application absent	Auto-applier dashboards
Perception–reality gap	Users feel more effective; measurement disagrees	−19% vs. +20% [6]
Cost-blind accounting	Subscription, review time, and reputation risk excluded	“Free” bulk applying
Fabrication externality	Embellished claims surface later (interview, reference, background check)	Ungrounded CV generation
Aggregate masking	Averages hide seniority/field segments where the tool fails	Cross-field conversion stats
Capability relabeling	Activity or template completion described as outcome value	Product capability claims

Table 3: Failure modes to test when evaluating AI career-tool claims. The taxonomy is an analytical checklist, not an estimate of category prevalence.

7 Where LLMs genuinely add value in the job search

The architecture follows a testable risk heuristic (Figure 2): assistance is easier to justify when errors are cheap to detect before action, and harder to justify when model output can directly trigger an irreversible outcome. Applied to this system:

Evaluate first: evidence-constrained drafting under operator review; research synthesis with surfaced citations (interview preparation, offer benchmarking); structured extraction checkable against a schema (paste-to-fill intake); deterministic quality signals (match reports, grounding gates) that inform rather than optimize; inspectable audit trails. The evaluation must include correction burden, false reassurance, and review time.

Avoid, or gate hard: unattended submission of applications; keyword optimization as a target rather than a signal; generation without a truth model; volume strategies that spend the candidate’s reputation; treating model confidence as calibrated probability.

In design terms: keep high-risk checks deterministic where possible; let no LLM authority cross the approval boundary; rank on observable signal and demote uncalibrated confidence; freeze baselines before interventions; and use the audit trail to make negative results diagnosable.

8 A checklist for evaluating AI career-tool claims

(1) Outcome data, or testimonial — and is the denominator (total applications) reported? (2) Net of all costs, including subscription, review time, and reputation risk? (3) What is n , and who selected the sample? (4) Stratified by seniority, field, and market, or aggregate? (5) Measured, or self-reported? (6) Who audits the number? (7) Is there a control arm — the same candidate, untailored or unassisted? (8) Would the vendor publish a negative result? (This paper prospectively specifies its measurements and commits to reporting failures so that item 8 is answered by construction.)




Figure 2: Risk heuristic. Model output may assist on the left, where errors are cheap to catch. Crossing the approval boundary makes remaining errors reputation-bearing. Supervision time belongs in the evaluation [19]; it is not a free safety layer.

9 Conclusion (provisional at v0.4)

ZenGrowth is a working verification-first career tool, not only a proposed architecture. It connects discovery, prioritization, evidence-grounded materials, interview preparation, offers, and later lifecycle stages while retaining reviewed evidence, deterministic checks for selected high-risk tokens, telemetry, and operator approval before dispatch. The repository and tests establish those useful capabilities. They do not yet quantify how smooth or fast the workflow is, the rate at which unsupported semantic claims survive, whether repeated scores are stable, or whether tailored materials improve callbacks. Those are prospective tests, not conclusions. Layered rewrite defenses with measured synthetic rates already exist [15]; what remains to be shown here is the operational and outcome value of a live, denominator-complete field deployment.

The novelty claim is correspondingly narrow. Prior systems already generate tailored resumes, retrieve from a career vault with provenance, and stack deterministic rewrite defenses. ZenGrowth may contribute a longitudinal field account connecting write-time controls, telemetry, correction burden, and downstream outcomes. It will publish complete denominators. Whether that contribution survives measurement is deliberately left open. At v0.4, the strongest result is a functioning, integrated workflow with tested controls, an explicit capability and outcome evidence ladder, and a falsifiable protocol for measuring its correctness and effectiveness.

Planned revisions. v1: execute and freeze the Table 2 measurements, publish the claim-level adjudication rubric, and report all aggregate results including failures. v2: add tracked-campaign outcome data, an executable replay fixture, and the final venue decision.

Reproducibility

ZenGrowth’s measurement protocol is documented in the repository [28]. The deterministic eval gates that witness architectural controls run in public CI. Dated aggregate baselines and

their derivation queries will be published with v1; underlying personal career data will not be published (see Limitations).

Acknowledgments

LLM-based coding and research tools assisted drafting, literature discovery, and implementation inside the ZenGrowth repository. The author reviewed the manuscript and accepts responsibility for its claims; empirical claims marked unmeasured remain unresolved by design. The author is the system’s sole operator and beneficiary, which is both the evidence base and a declared conflict. This work received no external funding.

References

- [1] Kumar Abhinav. Career-aware resume tailoring via multi-source retrieval-augmented generation with provenance tracking: A case study. <https://arxiv.org/abs/2605.05257>, 2026.
- [2] AIHawk contributors. AIHawk: Jobs applier AI agent. https://github.com/feder-cr/jobs_applier_ai_agent_aihawk, 2026. Open-source automated application agent; repository archived 2026-05-17. Accessed 2026-07.
- [3] Ada Aka, Emil Palikot, Ali Ansari, and Nima Yazdani. Better together: Quantifying the benefits of AI-assisted recruitment. <https://arxiv.org/abs/2507.08029>, 2026. Version 2, revised 2026-08-07.
- [4] Anthropic. What’s new in Claude Sonnet 5. <https://platform.claude.com/docs/en/models/sonnet-5/whats-new-sonnet-5>, 2026. Sampling-parameter and tokenizer behavior. Accessed 2026-08-25.
- [5] Lena Armstrong, Abbey Liu, Stephen MacNeil, and Danae Metaxa. The silicon ceiling: Auditing GPT’s race and gender biases in hiring. <https://arxiv.org/abs/2405.04412>, 2024.
- [6] Joel Becker, Nate Rush, Elizabeth Barnes, and David Rein. Measuring the impact of early-2025 AI on experienced open-source developer productivity. <https://arxiv.org/abs/2507.09089>, 2025.
- [7] Joel Becker, Nate Rush, Tom Cunningham, David Rein, and Khalid Mahamud. We are changing our developer productivity experiment design. <https://metr.org/blog/2026-02-24-uplift-update/>, 2026. Accessed 2026-07.
- [8] Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg more employable than Lakisha and Jamal? a field experiment on labor market discrimination. *American Economic Review*, 94(4): 991–1013, 2004. doi: 10.1257/0002828042002561.
- [9] Erik Brynjolfsson, Danielle Li, and Lindsey R. Raymond. Generative AI at work. Working Paper 31161, National Bureau of Economic Research, 2023.
- [10] Bhavya Chopra, Meng Chen, Rebecca Dang, Chanbin Park, Shreya Shankar, Sepanta Zeighami, Bjoern Hartmann, and Aditya Parameswaran. Who’s keeping score? interactive steering of LLM-powered scoring with Attune. <https://arxiv.org/abs/2608.14948>, 2026. To appear at ACM UIST 2026.
- [11] Shipi Dhanorkar, Samir Passi, and Mihaela Vorvoreanu. Human oversight of agentic systems in practice: Examining the oversight work, challenges, and heuristics of developers using software agents. <https://arxiv.org/abs/2606.05391>, 2026.
- [12] Ansel Kaplan Erol, Seohee Yoon, Keenan Hom, and Xisheng Zhang. Synapse: Evolving job-person fit with explainable two-phase retrieval and LLM-guided genetic resume optimization. <https://arxiv.org/abs/2604.02539>, 2026.
- [13] Zhenyu Gao, Wenxi Jiang, and Yutong Yan. Can LLMs hire fairly? racial bias in resume screening. <https://arxiv.org/abs/2606.28978>, 2026.
- [14] Serena A. Hoffstedde, Machiko Hirota, Akshara Nadayanur Sathis Kanna, Rihito Kotani, Ujwal Kumar, Gabriele Trovato, and Phan Xuan Tan. Gender bias in LLM hiring decisions: Evidence from

- a japanese context and evaluation of mitigation strategies. <https://arxiv.org/abs/2606.18649>, 2026.
- [15] Shashank Indukuri and Adarsh Agrawal. Grounded optimization: A layered engineering framework for reducing LLM hallucination in automated personal document rewriting. <https://arxiv.org/abs/2607.01457>, 2026.
 - [16] JobCopilot. Jobcopilot: Automated job applications. <https://jobcopilot.com/>, 2026. Automated-application capability and outcome marketing claims. Accessed 2026-07.
 - [17] Jobscan. Jobscan: Resume scanner and match rate. <https://www.jobscan.co/resume-matcher>, 2026. Product-defined Match Rate and 75% recommendation. Accessed 2026-07.
 - [18] Keyu Li, Junhao Shi, Yang Xiao, Mohan Jiang, Jie Sun, Yunze Wu, Dayuan Fu, Shijie Shi, Xiaojie Cai, Tianze Xu, Weiye Si, Wenjie Li, Dequan Li, and Pengfei Liu. AgencyBench: Benchmarking the frontiers of autonomous agents in 1m-token real-world contexts. <https://arxiv.org/abs/2601.11044>, 2026. Accepted at ACL 2026.
 - [19] Weiwen Liu, Jiarui Qin, Xu Huang, Xingshan Zeng, Yunjia Xi, Jianghao Lin, Chuhan Wu, Yasheng Wang, Lifeng Shang, Ruiming Tang, Defu Lian, Yong Yu, and Weinan Zhang. Position: The real barrier to LLM agent usability is agentic ROI. <https://arxiv.org/abs/2505.17767>, 2025.
 - [20] Yi Nian, Aojie Yuan, Haiyue Zhang, Jiate Li, Li Li, Xiyang Hu, Hua Wei, Xiongye Xiao, Chaowei Xiao, and Yue Zhao. Auditable agents. <https://arxiv.org/abs/2604.05485>, 2026.
 - [21] Sida Peng, Eirini Kalliamvakou, Peter Cihon, and Mert Demirer. The impact of AI on developer productivity: Evidence from GitHub Copilot. <https://arxiv.org/abs/2302.06590>, 2023.
 - [22] Bryan Chen Zhengyu Tan, Shaun Khoo, Bich Ngoc Doan, Zhengyuan Liu, Nancy F. Chen, and Roy Ka-Wei Lee. Small changes, big impact: Demographic bias in LLM-based hiring through subtle sociocultural markers in anonymised resumes. <https://arxiv.org/abs/2603.05189>, 2026.
 - [23] Teal Labs. Teal: Job application tracker and resume match score. <https://www.tealhq.com/tool/resume-job-description-match>, 2026. Resume-to-job Match Score product page. Accessed 2026-07.
 - [24] Kyra Wilson and Aylin Caliskan. Gender, race, and intersectional bias in resume screening via language model retrieval. <https://arxiv.org/abs/2407.20371>, 2024.
 - [25] Kyra Wilson, Mattea Sim, Anna-Maria Gueorguieva, and Aylin Caliskan. No thoughts just AI: Biased LLM hiring recommendations alter human decision making and limit human autonomy. <https://arxiv.org/abs/2509.04404>, 2025. Published in the Proceedings of the 2025 AAAI/ACM Conference on AI, Ethics, and Society.
 - [26] Xiao Yu, Ruize Xu, Chengyuan Xue, Junyu Chen, Matthew So, Shijun Ma, Bo Liu, Xiangye Liang, and Zhou Yu. ConFit v3: Improving resume-job matching with LLM-based re-ranking. <https://arxiv.org/abs/2605.09760>, 2026.
 - [27] Cesare Zavattari, Alessandro Tommasi, and Giuseppe Prencipe. One human, n agents: Audit-budget allocation for LLM agent fleets under miscalibrated, correlated confidence. <https://arxiv.org/abs/2607.28317>, 2026.
 - [28] ZenGrowth. Zengrowth academic publication plan and measurement protocol. <https://github.com/zenouz-ai/zengrowth/blob/main/docs/ACADEMIC-PUBLICATION-PLAN.md>, 2026. Prospective procedures; dated aggregate baselines planned for v1. Accessed 2026-08.
 - [29] Saurabh Bhausaheb Zinjad, Amrita Bhattacharjee, Amey Bhilegaonkar, and Huan Liu. ResumeFlow: An LLM-facilitated pipeline for personalized resume generation and refinement. <https://arxiv.org/abs/2402.06221>, 2024.